



ALLIANCE
UNIVERSITY
Private University established in Karnataka State by Act No.34 of year 2010
Recognized by the University Grants Commission (UGC), New Delhi

NAAC
GRADE A+
ACCREDITED UNIVERSITY

Bachelor of Computer Applications

Semester – V

COURSE : Deep Learning

Project : 04

Title: Ethical Considerations in Deep Learning

Submitted By

Name: Laishram Ritikumar Singh

Reg. No.:2411021240040

Date:04-09-2026

Submitted To:

Faculty Name: Nagasri P.

Department of Computer Application
Alliance University
Chandapura - Anekal Main Road, Anekal
Bengaluru – 562106

Ethical Considerations in Deep Learning

Introduction

Deep learning has grown into one of the most influential technologies in computing today. It underlies applications such as facial recognition, medical diagnostics, credit-risk scoring, personalized recommendations, autonomous vehicles, and even tools used to assess risk within the justice system. Because these systems now shape decisions that touch people's employment, liberty, health, and safety, it becomes critical to closely examine the ethical dimensions of how they are built, trained, and rolled out.

In contrast to conventional software built on explicit rules, deep learning systems derive their behavior from the data used to train them. As a result, the quality, fairness, and safety of their outputs depend heavily on the data and design decisions behind the model. Whenever the underlying data is incomplete, skewed, or reflects historical prejudice—or when models are put into use without sufficient monitoring—the resulting damage can be severe and hard to undo. This report explores three prominent ethical concerns in deep learning, reviews a documented real-world case where one of these concerns led to tangible harm, and outlines practical steps that can help mitigate such risks in the future.

Task 1: Three Ethical Issues in Deep Learning

Deep learning brings up numerous ethical concerns, but three stand out for their direct impact on fairness, individual liberty, and public confidence: bias and discrimination, privacy intrusion, and limited transparency. Each of these is discussed in more depth below.

1. Bias and Discrimination

Deep learning models identify patterns from historical records. When that history carries forward existing social, economic, or cultural inequities—for instance, when particular racial, gender, or socioeconomic groups were previously underrepresented, mischaracterized, or unfairly treated—the model absorbs and repeats these same patterns. At times, models go further than simple repetition and actually intensify the bias, since the training process rewards the model for aligning with majority trends in the data, regardless of whether those trends are equitable.

Bias can be introduced into a deep learning system at several points along the way: while gathering data (if some groups are inadequately represented), while labeling data (if the people annotating it carry their own biases), while selecting features (if variables used as proxies correlate with protected traits), and while training the model (if fairness is not built into the objective function). Figure 1 shows how bias can be introduced at the very first stage of the pipeline and then

strengthened over time through a feedback loop, as biased outcomes generate additional biased data.

Figure 1: How Bias Enters and Is Reinforced in a Deep Learning Pipeline




Figure 1: The pathway by which bias enters and is reinforced across stages of a deep learning pipeline.

These consequences become particularly severe in high-stakes areas like recruitment, lending, healthcare, and the criminal justice system, where a flawed prediction can determine whether an individual is hired, granted a loan, given appropriate medical care, or released on parole.

2. Privacy Violations

Deep learning systems, particularly those operating at large scale, are usually trained using enormous datasets that can contain photographs, audio recordings, text conversations, browsing activity, biometric identifiers, or medical histories. A large portion of this information is personal and sensitive in nature. Gathering, storing, and using such data without explicit consent, sufficient protections, or a valid justification poses a substantial threat to individual privacy.

Privacy concerns in deep learning show up in multiple ways. First, the training data may be gathered without genuine consent—for example, through automated scraping of images or text found online. Second, trained models occasionally retain and expose pieces of their training data, an issue referred to as model inversion or membership inference, in which an attacker can figure out whether a particular person's information was part of the training set. Third, after deployment, the outputs a model produces—such as a health-risk estimate or a facial-recognition match—can themselves disclose sensitive details about a person, even without the original data ever being directly exposed.

3. Lack of Transparency (The Black-Box Problem)

A great many deep learning models, especially deep neural networks composed of numerous hidden layers, are frequently labeled "black boxes," since it is remarkably hard—even for the engineers who designed them—to pinpoint exactly why a model arrived at a given output. A model may deliver correct predictions the majority of the time, yet still leave no one able to clearly identify which features drove the outcome for any specific case.




Figure 2: Inside a typical deep neural network, the conversion of input into output takes place through numerous hidden layers that remain opaque to human interpretation.

This absence of transparency gives rise to several further complications. It becomes harder to notice when a model is drawing on inappropriate or biased cues. It becomes harder for people affected by a decision to understand or push back against it. And it becomes harder for regulators, auditors, or courts to hold organizations to account, since a clear explanation is frequently unavailable for evaluation. Transparency is thus tightly connected to both fairness and accountability: in its absence, bias can remain concealed inside a system for extended periods.

Task 2: Real-World Example – The COMPAS Recidivism Algorithm

A thoroughly documented real-world case that brings together all three issues discussed above is that of COMPAS (Correctional Offender Management Profiling for Alternative Sanctions), a proprietary risk-assessment tool built by Northpointe (later renamed Equivant). COMPAS applies a statistical and machine-learning-driven model to produce a "risk score" reflecting the likelihood that a criminal defendant will reoffend. Beginning in the early 2000s, courts across several U.S. states, Wisconsin among them, have relied on COMPAS scores to help guide decisions on bail, sentencing, and parole.

What Happened

In May 2016, the investigative journalism outlet ProPublica released an extensive statistical review of COMPAS risk scores covering more than 7,000 defendants arrested in Broward County,

Florida. The investigation set predicted risk scores against actual reoffense outcomes tracked over two years. It uncovered a marked racial disparity in the algorithm's errors: Black defendants who ultimately did not reoffend were nearly twice as likely to be wrongly flagged as "high risk" compared with white defendants who also did not reoffend. On the other hand, white defendants who did go on to reoffend were more frequently mislabeled as "low risk" than Black defendants who reoffended. Although overall accuracy was comparable across racial groups, the distribution of errors was not, meaning the burden of the algorithm's mistakes fell disproportionately on Black defendants.

The inner workings and scoring methodology of the tool were never made fully public, since COMPAS is proprietary software that its developer guarded as a trade secret. This issue became central to *Loomis v. Wisconsin* (2016), in which defendant Eric Loomis contested the use of his COMPAS score at sentencing, arguing he had no way to review or dispute how the score had been derived. The Wisconsin Supreme Court ultimately permitted continued use of COMPAS but advised judges to apply the scores cautiously, and the U.S. Supreme Court subsequently declined to review the matter further, leaving the underlying transparency issues largely unaddressed.

Linking the Incident to Task 1

This case offers a vivid real-world demonstration of the three ethical issues introduced earlier, and it illustrates how they interact rather than operate independently.

- **Bias and Discrimination:** The algorithm generated consistently unequal error rates between racial groups, even though race itself was not a direct input, because proxy factors—such as neighborhood, employment record, and family arrest history—correlated with race and embedded historical inequity into the model's predictions.
- **Lack of Transparency:** Since COMPAS operated as a closed, proprietary system, defendants, their attorneys, and the courts had no way to fully examine how any given score was produced, making it exceedingly difficult to detect or formally contest the source of bias before ProPublica's independent statistical review.
- **Privacy:** Sensitive personal, behavioral, and demographic information about defendants was fed into the system to produce scores that meaningfully shaped decisions on bail, sentencing, and parole, frequently without defendants having a clear understanding of what data was used or how it influenced the result.

The COMPAS case is frequently referenced in AI ethics scholarship precisely because it shows, backed by concrete statistical evidence and an actual legal challenge, how bias, opacity, and privacy concerns can converge within a single deployed deep-learning-adjacent system to inflict genuine harm on people's lives and freedom.

Task 3: Measures to Address the Identified Issues

Tackling the ethical issues outlined above calls for coordinated action across several levels: technical safeguards embedded within the model itself, organizational policies governing data and model management, and external oversight mechanisms that hold developers and deployers accountable. Figure 3 presents these measures as three supporting pillars of a responsible AI system.

Figure 3: A Framework for Responsible Deep Learning Systems




Figure 3: A framework illustrating how technical, organizational, and oversight measures jointly sustain a trustworthy AI system.

Measures to Reduce Bias and Discrimination

- Rely on diverse and representative training datasets that properly capture all relevant demographic groups, rather than depending on convenience samples.
- Carry out periodic fairness audits that assess model error rates and outcomes separately by race, gender, age, and other protected characteristics, similar to the review ProPublica conducted on COMPAS.
- Employ fairness-aware machine learning approaches—such as re-weighting training samples, adversarial de-biasing, or post-processing calibration—to narrow disparities in outcomes.
- Identify and closely scrutinize proxy variables (like zip code or neighborhood) that might indirectly capture protected characteristics such as race or socioeconomic background.

Measures to Protect Privacy

- Follow the principle of data minimization: gather only what is genuinely required for the task, and secure clear, informed consent from the individuals whose data is being used.
- Adopt privacy-preserving methods such as anonymization, differential privacy, and encryption for data both at rest and in transit.
- Adhere to relevant data protection regulations, such as the General Data Protection Regulation (GDPR) or comparable national laws, which mandate lawful grounds for processing personal data and grant individuals rights over their own information.
- Restrict how long data is retained and clearly specify who may access sensitive data and for what purposes.

Measures to Improve Transparency and Accountability

- Apply Explainable AI (XAI) methods, including feature-importance analysis and model-agnostic explanation tools, so that developers and affected individuals alike can understand the reasoning behind a given decision.
- Permit independent, third-party audits of high-stakes systems—especially in criminal justice, healthcare, hiring, and finance—rather than depending solely on internal reviews.
- Set up clear accountability structures that specify responsibility when an AI system causes harm, including avenues for individuals to appeal or contest automated decisions.
- Release model documentation, such as "model cards" or "datasheets for datasets," detailing a system's intended use, its limitations, and any known performance gaps across groups.

Conclusion

Deep learning delivers substantial benefits, yet as the COMPAS case makes clear, its ethical risks are far from hypothetical. Bias, privacy intrusion, and limited transparency are interrelated issues that can combine to produce serious, tangible harm, particularly for groups that are already vulnerable. Meeting these challenges takes more than good intentions: it demands carefully curated and audited data, privacy-conscious engineering, models that can be explained and audited, and well-defined accountability structures. When technical safeguards are paired with robust organizational policy and independent oversight, deep learning systems can be deployed in ways that are fairer, safer, and more responsible.